

SKYGUARD AI

6-Layer Fusion for AWS Anomaly Detection

Physics • ML • Forecasting • Spatial Consensus — WMO-L0-L5 aligned

Problem Statement: AI/ML-Based Intelligent Anomaly Detection for Automatic Weather Stations (AWS)

Theme: Smart Automation & Observational Networks | **Category:** Software (with Edge-AI hardware)

Parameters: Temperature °C • Atmospheric Pressure hPa • Relative Humidity %

Team ID: _____ | **Team Name:** _____ | **Mentor:** _____

Solution — 6-Layer Fusion mirroring WMO L0-L5



Layer Roles

- **L1 Physics** — range, dew-point (Magnus), barometric altitude, rate-of-change. Deterministic.
- **L2 ML** — Extended IsoForest + SHAP + LSTM-AE + Mann-Kendall drift + gap detection.
- **L3 Frozen** — $\text{std} < 0.01 / d^2x/dt^2 \approx 0$; all-3-frozen → logger failure.
- **L4 Forecasting** — Chronos-2 (120M, multivariate) corridor breach.
- **L5 Spatial** — robust $Z = |x - \text{median}| / (1.4826 \cdot \text{MAD})$; MSL-normalize pressure.
- **L6 Fusion** — 2x2 matrix → verdict + confidence + fault type + action.

Why 6 layers beat any single model

IsoForest misses frozen → L3 catches • misses slow drift → L2 trend + L4 residual • threshold-only misses multivariate combos → L1 dew-point + L2 • any detector false-positives a heatwave → L5 agreement → REGIONAL_EVENT • black-box alert → SHAP + per-fault natural language.

The Fusion Matrix

	Neighbours AGREE ($Z < 3.1$)	Neighbours DISAGREE ($Z \geq 3.1$)
Model flagged (L1-L4)	REGIONAL_EVENT real weather — alert, not fault — Medium	SENSOR_FAULT sensor is broken — High
No model flag	HEALTHY all clear — no action	CHECK_HEALTH subtle drift — schedule maintenance — Low

Novelty

Full integration of edge-deployed ML (ESP32) + physics + foundation-forecast + spatial fusion + self-healing health loop — **not published as a complete system**. Components individually validated by ECMWF (operational 2023) / Met Éann (99.6%, 2026) / IEEE 2026 edge-IF.

Self-Healing Loop (grand-challenge answer)

Composite sensor-health score + predictive time-to-failure + auto-imputation → degraded sensor → downstream feeds upweight imputed values until fixed → **trustworthy data under all conditions**.

Concrete Thresholds (all sourced)

Check	Threshold	Source
Temp rate (fault)	$ \Delta T > 19.4 \text{ }^{\circ}\text{C/hr}$	NOAA MADIS (35°F/hr)
Temp rate (alert)	$ \Delta T > 8 \text{ }^{\circ}\text{C/hr}$	Tightened watch level
Pressure rate (fault)	$ \Delta P > 15 \text{ hPa/hr}$	NOAA MADIS sea-level
Pressure rate (alert)	$ \Delta P > 3 \text{ hPa/hr}$	Cyclone-class (weather, not fault)
Humidity rate (fault)	$ \Delta RH > 50 \text{ \% /hr}$	NOAA MADIS
Dew-point constraint	$T_{\text{dew}} \leq T_{\text{air}} + 0.5^{\circ}\text{C}$	Magnus / Alduchov-Eskridge; MADIS
Barometric altitude	$ P - P_{\text{exp}} > 30 \text{ hPa}$	Barometric formula
Range bounds (India)	$T -40/+55^{\circ}\text{C}$, RH 0-100%, P 300-1084	WMO records, BoM
Climatological	monthly mean $\pm 5\sigma$	WMO guideline / Lanzante 1996
Spatial robust-Z	$Z \geq 3.1 \rightarrow \text{disagree}$	Met Norway TITAN default
Frozen sensor	std < 0.01 over ≥ 10 readings	OK Mesonet persistence test
Forecast breach	actual $\notin [p5, p95]$ or severity > 0.8	Chronos-2 corridor
Drift (Mann-Kendall)	$p < 0.05 + \text{Theil-Sen slope}$	Standard trend test
ECMWF blacklist	$PGE > 0.75$	ECMWF operational

Tuning: sweep each threshold on injected-anomaly benchmark; maximize **F1 / F2 / PR-AUC** (not accuracy — anomalies $< 5\%$ of data). Per-station monthly $\pm 5\sigma$ replaces global range in production.

Explainability (XAI)

- **SHAP TreeExplainer** on IsoForest — `check_additivity=False`; explains path length (sign gotcha documented).
- Blamed sensor = feature with largest $|\text{SHAP}|$.
- Confidence tier from #layers agreeing + spatial.

Production-Grade Maturity Ladder

Stage	Scope	Stack
Idea (now)	Slides + architecture; injected anomalies on ERA5	Markdown plan
Pilot	1 region, ~50 stations, real IMD API, live dashboard	Docker Compose; A/B vs deterministic QC
Production	675+ IMD, 24/7, NWP O-B virtual buddy, audit logs	K8s autoscale + BullMQ; Grafana; WMO traceability

MLOps (the real gap)

- **Model versioning:** DVC/MLflow per station cluster
- **Drift monitoring:** PSI > 0.2 → retrain (Kataria 2026, F1=0.91)
- **A/B thresholds:** deterministic + ML in parallel
- **Human-in-loop:** AI proposes blacklist, analyst dispositions (ECMWF model)
- **Audit log:** every QC decision traceable (WMO-No. 8 / WIGOS)
- **Fallback:** if ML down → L1+L3+L5 deterministic only

Risks & Mitigations

Failure mode	Mitigation
Isolated / terrain / frontal	Multi-layer + NWP virtual buddy + lapse-rate + inversion detector
IsoForest ghost clusters / drift	Extended IsoForest + LSTM-AE + PSI retrain + de-seasonalize
ML flags real extreme weather	Spatial consensus + climatological bounds + exclude REGIONAL_EVENT from retrain
Class imbalance	PR-AUC / F2 / precision-at-recall, not accuracy
Labelled-data scarcity	Injected benchmark + semi-auto labels from ops warnings (ECMWF)
ESP32 constraints	int8 + OTA rollback + ADC calibration
Regulatory acceptance	Explainability + human-in-loop + deterministic fallback + audit (WMO Oct 2025)
Telemetry integrity (WMO gap)	HMAC-signed readings + tamper classifier

Impact

- **Direct:** IMD 675+ AWS → cleaner data → better forecasts, disaster warnings, aviation safety.
- **Scale:** every national met service (NOAA 20,000+, ECMWF, DWD, BoM, Met Office).
- **3-parameter scope** → globally applicable to any T/P/RH AWS.

Benefits

Type	Benefit
Social	Lives saved via better disaster warnings
Economic	Fewer wrong forecasts; predictive maintenance lowers cost
Environmental	72% energy saving; longer field life; less e-waste
Scientific	Trustworthy climate-record inputs; traceable QC

Future Work Progression

- Probabilistic Constant Value Test (AMT 2023) for frozen detection
- Conformal prediction for calibrated coverage intervals
- NWP O-B “virtual buddy” for isolated stations
- Federated learning across stations (privacy-preserving scale)
- Physics-Informed Neural Network hybrid

User Experience

Every alert = verdict + confidence (HIGH/MED/LOW) + blamed sensor (SHAP) + plain-English reason + recommended action + imputed value + source.

Dashboard

- Leaflet map — station colour = health
- Live alert feed — severity colour-coded, per-layer expandable
- SHAP bar chart per flagged reading
- Chronos-2 forecast corridor (p5–p95) with actual marked
- Sensor-health panel + predicted time-to-failure

Expected Outputs (problem statement)

Output	Delivered by
Real-time alerts	L6 → WebSocket
Severity + confidence	L6 tier + numeric
Root-cause classification	Fault-type priority (Frozen>Physics>Forecast>ML)
Visualization dashboard	React + Leaflet + Recharts
Sensor health status	Composite score [0,1] + tier
Corrected data (optional)	Imputation ensemble (neighbour → Kalman → Chronos → NWP analog)

1. WMO-No. 8, Guide to Instruments & Methods of Observation (CIMO), 2024 ed. — <https://library.wmo.int/viewer/68695/>
2. NOAA MADIS surface QC notes — https://madis.ncep.noaa.gov/madis_sfc_qc_notes.shtml
3. ECMWF, Dahoui et al. (2023), ML detection & classification of observation anomalies, Newsletter #174 — <https://www.ecmwf.int/en/newsletter/174/>
4. Spohn et al. (2026), Autoencoders for met QC, Met Éann, **99.6% accuracy**, Environmental Data Science — <https://www.cambridge.org/core/journals/environmental-data-science/article/machine-learning-approach-using-autoencoders-to-perform-quality-control-on-meteorological-data/4576781508080877E36C0CA6612E5590>
5. Edge-Optimized Isolation Forest, IEEE ICISCOIS 2026 (**<10ms, <200KB**) — <https://doi.org/10.1109/iciscois62701.2026.11447816>
6. Kataria et al. (2026), 6-layer MLOps pipeline, **F1=0.91**, IEEE MENACOMM — <https://doi.org/10.1109/menacomm69507.2026.11532803>
7. Met Norway TITAN QC framework (Løpez et al. 2020) — <https://asr.copernicus.org/articles/17/153/2020/> | ROVE engine — <https://github.com/metno/rove>
8. Chronos-2, Amazon (arXiv:2510.15821, Oct 2025, 120M params, multivariate) — <https://huggingface.co/amazon/chronos-2>
9. SHAP TreeExplainer on IsolationForest (PR #784) — <https://github.com/shap/shap/pull/784>
10. alibi-detect v0.13.0 (Seldon) — <https://github.com/SeldonIO/alibi-detect>
11. Fiebrich & Crawford (2001), Oklahoma Mesonet QA, Bull. Amer. Meteor. Soc. — [https://doi.org/10.1175/1520-0477\(2001\)082](https://doi.org/10.1175/1520-0477(2001)082)
12. ESP32-S3 Datasheet (Espressif) — https://documentation.espressif.com/esp32-s3_datasheet_en.html
13. Open-Meteo Archive API (ERA5, free, no key) — <https://open-meteo.com/en/docs/historical-weather-api>
14. NOAA Integrated Surface Database (ISD) — <https://www.ncei.noaa.gov/products/land-based-station/integrated-surface-database>
15. De Bruijn et al. (2016), injected-fault sensor benchmark, 5.78M points — <https://doi.org/10.5220/0005637901850195>
16. WMO Congress (Oct 2025), AI complements not replaces physics, transparency & traceability — <https://wmo.int/themes/artificial-intelligence>
17. SIH 2026 Idea Presentation Format (6-slide PDF) — <https://www.scribd.com/document/1075380264/SIH2026-IDEA-Presentation-Format>